

# Bayesian Learning

## Lecture 9 - HMC, Stan and Variational Inference (VI)

Bertil Wegmann

Department of Computer and Information Science  
Linköping University



# Lecture overview

- Hamiltonian Monte Carlo
- Stan
- Variational Inference

# Hamiltonian Monte Carlo

- When  $\theta = (\theta_1, \dots, \theta_p)$  is **high-dimensional**,  $p(\theta|y)$  usually located in some subregion of  $\mathbb{R}^p$  with complicated geometry.
- MH: hard to find good proposal distribution  $q(\cdot|\theta^{(i-1)})$ .
- MH: use very small step sizes otherwise too many rejections.
- **Hamiltonian Monte Carlo (HMC)**:
  - ▶ distant proposals **and**
  - ▶ high acceptance probabilities.
- HMC: add extra **momentum** parameters  $\phi = (\phi_1, \dots, \phi_p)$  and sample from

$$p(\theta, \phi|y) = p(\theta|y) p(\phi)$$

# Hamiltonian Monte Carlo

- Physics: **Hamiltonian** system  $H(\theta, \phi) = U(\theta) + K(\phi)$ , where  $U$  is the **potential energy** and  $K$  is the **kinetic energy**.
- **Hamiltonian Dynamics**

$$\begin{aligned}\frac{d\theta_i}{dt} &= \frac{\partial H}{\partial \phi_i} = \frac{\partial K}{\partial \phi_i}, \\ \frac{d\phi_i}{dt} &= -\frac{\partial H}{\partial \theta_i} = -\frac{\partial U}{\partial \theta_i}\end{aligned}$$

- Hockey puck sliding over a friction-less surface: [illustration](#).
- Use  $U(\theta) = -\log[p(\theta)p(y|\theta)]$ .
- Use  $\phi \sim N(0, M)$  where  $M$  is the mass matrix and

$$K(\phi) = -\log[p(\phi)] = \frac{1}{2}\phi^T M^{-1}\phi + \text{const}$$

- If we could propose  $\theta$  in continuous time (spoiler: we can't), the acceptance probability would be one.

# Hamiltonian Monte Carlo

## ■ Hamiltonian Dynamics

$$\begin{aligned}\frac{d\theta_i}{dt} &= [M^{-1}\phi]_i, \\ \frac{d\phi_i}{dt} &= \frac{\partial \log p(\theta|y)}{\partial \theta_i}\end{aligned}$$

which can be simulated using the **leapfrog algorithm**

$$\begin{aligned}\phi_i\left(t + \frac{\varepsilon}{2}\right) &= \phi_i(t) + \frac{\varepsilon}{2} \frac{\partial \log p(\theta(t)|y)}{\partial \theta_i}, \\ \theta(t + \varepsilon) &= \theta(t) + \varepsilon M^{-1}\phi(t), \\ \phi_i(t + \varepsilon) &= \phi_i\left(t + \frac{\varepsilon}{2}\right) + \frac{\varepsilon}{2} \frac{\partial \log p(\theta(t)|y)}{\partial \theta_i},\end{aligned}$$

where  $\varepsilon$  is the step size.

■ **Discretization**  $\Rightarrow$  acceptance probability drops with  $\varepsilon$ .

# The Hamiltonian Monte Carlo algorithm

■ Initialize  $\theta^{(0)}$  and iterate for  $i = 1, 2, \dots$

- 1 Sample the starting **momentum**  $\phi_s \sim N(0, M)$
- 2 Simulate new values for  $(\theta_p, \phi_p)$  by iterating the **leapfrog algorithm**  $L$  times, starting in  $(\theta^{(i-1)}, \phi_s)$ .
- 3 Compute the **acceptance probability**

$$\alpha = \min \left( 1, \frac{p(y|\theta_p)p(\theta_p)}{p(y|\theta^{(i-1)})p(\theta^{(i-1)})} \frac{p(\phi_p)}{p(\phi_s)} \right)$$

- 4 With probability  $\alpha$  set  $\theta^{(i)} = \theta_p$  and  $\phi^{(i)} = \phi_p$  otherwise.

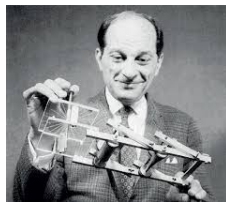
■ **Tuning parameters:** 1. **stepsize**  $\varepsilon$ , 2. **number of leapfrog iterations**  $L$  and 3. **mass matrix**  $M$ . **No U-turn.**

# Stan

- **Stan** is a probabilistic programming language based on HMC.
- Allows for Bayesian inference in many models with automatic implementation of the MCMC sampler.
- Named after Stanislaw Ulam (1909-1984), co-inventor of the Monte Carlo algorithm.
- Written in C++ but can be run from R using the package `rstan`




Stan logo



Stanislaw Ulam

# Stan - toy example: three plants

- Three plants were observed for four months, measuring the number of flowers





# Stan Model 1: iid normal

$$y_i \stackrel{iid}{\sim} N(\mu, \sigma^2)$$

```
library(rstan)
y=c(4,5,6,4,0,2,5,3,8,6,10,8)
N=length(y)

StanModel = '
data {
  int<lower=0> N; // Number of observations
  int<lower=0> y[N]; // Number of flowers
}
parameters {
  real mu;
  real<lower=0> sigma2;
}
model {
  mu ~ normal(0,100); // Normal with mean 0, st.dev. 100
  sigma2 ~ scaled_inv_chi_square(1,2); // Scaled-inv-chi2 with nu 1, sigma 2
  for(i in 1:N){
    y[i] ~ normal(mu,sqrt(sigma2));
  }
},'
```

# Stan Model 2: multilevel normal

$$y_{t,p} \sim N\left(\mu_p, \sigma_p^2\right), \mu_p \sim N\left(\mu, \sigma^2\right)$$

```
StanModel <- '  
data {  
  int<lower=0> N; // Number of observations  
  int<lower=0> y[N]; // Number of flowers  
  int<lower=0> P; // Number of plants  
}  
transformed data {  
  int<lower=0> M; // Number of months  
  M = N / P;  
}  
parameters {  
  real mu;  
  real<lower=0> sigma2;  
  real mup[P];  
  real sigmap2[P];  
}  
model {  
  mu ~ normal(0,100); // Normal with mean 0, st.dev. 100  
  sigma2 ~ scaled_inv_chi_square(1,2); // Scaled-inv-chi2 with nu 1, sigma 2  
  for(p in 1:P){  
    mup[p] ~ normal(mu,sqrt(sigma2));  
    for(m in 1:M) {  
      y[M*(p-1)+m] ~ normal(mup[p],sqrt(sigmap2[p]));  
    }  
  }  
}
```

# Stan Model 3: multilevel Poisson

$$y_{t,p} \sim \text{Poisson}(\mu_p), \mu_p \sim \text{logN}(\mu, \sigma^2)$$

```
StanModel <- '
data {
  int<lower=0> N; // Number of observations
  int<lower=0> y[N]; // Number of flowers
  int<lower=0> P; // Number of plants
}
transformed data {
  int<lower=0> M; // Number of months
  M = N / P;
}
parameters {
  real mu;
  real<lower=0> sigma2;
  real mup[P];
}
model {
  mu ~ normal(0,100); // Normal with mean 0, st.dev. 100
  sigma2 ~ scaled_inv_chi_square(1,2); // Scaled-inv-chi2 with nu 1, sigma 2
  for(p in 1:P){
    mup[p] ~ lognormal(mu,sqrt(sigma2)); // Log-normal
    for(m in 1:M) {
      y[M*(p-1)+m] ~ poisson(mup[p]); // Poisson
    }
  }
}'
```

# Stan: fit model and analyze output

```
data <- list(N=N, y=y, P=P)
warmup <- 1000
niter <- 2000
fit <- stan(model_code=StanModel,data=data, warmup=warmup,iter=niter,chains=4)

# Print the fitted model
print(fit,digits_summary=3)

# Extract posterior samples
postDraws <- extract(fit)

# Do traceplots of the first chain
par(mfrow = c(1,1))
plot(postDraws$mu[1:(niter-burnin)],type="l",ylab="mu",main="Traceplot")

# Do automatic traceplots of all chains
traceplot(fit)

# Bivariate posterior plots
pairs(fit)
```

# Stan - useful links

- [Getting started with RStan](#)
- [RStan vignette](#)
- [Stan Modeling Language User's Guide and Reference Manual](#)
- [Stan Case Studies](#)

# Variational Inference

- Let  $\theta = (\theta_1, \dots, \theta_p)$ . Approximate the posterior  $p(\theta|y)$  with a (simpler) distribution  $q(\theta)$ .
- Before: **Normal approximation** from optimization:  
 $q(\theta) = N[\tilde{\theta}, J_y^{-1}(\tilde{\theta})]$ .
- **Mean field Variational Inference (VI)**:  $q(\theta) = \prod_{i=1}^p q_i(\theta_i)$
- **Parametric VI**: Parametric family  $q_\lambda(\theta)$  with parameters  $\lambda$
- Find the  $q(\theta)$  that **minimizes the Kullback-Leibler distance** between the true posterior  $p$  and the approximation  $q$ :

$$KL(q, p) = \int q(\theta) \ln \frac{q(\theta)}{p(\theta|y)} d\theta = E_q \left[ \ln \frac{q(\theta)}{p(\theta|y)} \right].$$

# Mean field approximation

- **Mean field VI** is based on factorized approximation:

$$q(\theta) = \prod_{i=1}^p q_i(\theta_i)$$

- **No specific functional forms** are assumed for the  $q_i(\theta)$ .
- **Optimal densities** can be shown to satisfy:

$$q_j(\theta) \propto \exp(E_{-\theta_j} \ln p(y, \theta))$$

where  $E_{-\theta_j}(\cdot)$  is the expectation with respect to  $\prod_{k \neq j} q_k(\theta_k)$ .

# Mean field approximation - algorithm

- Initialize:  $q_2^*(\theta_2), \dots, q_p^*(\theta_p)$
- Repeat until convergence:
  - ▶  $q_1^*(\theta_1) \leftarrow \frac{\exp[E_{-\theta_1} \ln p(y, \theta)]}{\int \exp[E_{-\theta_1} \ln p(y, \theta)] d\theta_1}$
  - ▶  $\vdots$
  - ▶  $q_p^*(\theta_p) \leftarrow \frac{\exp[E_{-\theta_p} \ln p(y, \theta)]}{\int \exp[E_{-\theta_p} \ln p(y, \theta)] d\theta_p}$
- Note: no assumptions about parametric form of the  $q_i(\theta)$ .
- Optimal  $q_i(\theta)$  often **turn out** to be parametric (normal etc).
- Just update hyperparameters in the optimal densities.



# Mean field approximation - Normal model

- **Model:**  $X_i | \theta, \sigma^2 \stackrel{iid}{\sim} N(\theta, \sigma^2)$ .
- **Prior:**  $\theta \sim N(\mu_0, \tau_0^2)$  independent of  $\sigma^2 \sim \text{Inv} - \chi^2(\nu_0, \sigma_0^2)$ .
- **Mean-field approximation:**  $q(\theta, \sigma^2) = q_\theta(\theta) \cdot q_{\sigma^2}(\sigma^2)$ .
- Optimal densities

$$q_\theta^*(\theta) \propto \exp \left[ E_{q(\sigma^2)} \ln p(\theta, \sigma^2, \mathbf{x}) \right]$$
$$q_{\sigma^2}^*(\sigma^2) \propto \exp \left[ E_{q(\theta)} \ln p(\theta, \sigma^2, \mathbf{x}) \right]$$

# Normal model - VB algorithm

## ■ Variational density for $\sigma^2$

$$\sigma^2 \sim \text{Inv} - \chi^2 (\tilde{\nu}_n, \tilde{\sigma}_n^2)$$

where  $\tilde{\nu}_n = \nu_0 + n$  and  $\tilde{\sigma}_n^2 = \frac{\nu_0 \sigma_0^2 + \sum_{i=1}^n (x_i - \tilde{\mu}_n)^2 + n \cdot \tilde{\tau}_n^2}{\nu_0 + n}$

## ■ Variational density for $\theta$

$$\theta \sim N(\tilde{\mu}_n, \tilde{\tau}_n^2)$$

where

$$\tilde{\tau}_n^2 = \frac{1}{\frac{n}{\tilde{\sigma}_n^2} + \frac{1}{\tau_0^2}}$$


$$\tilde{\mu}_n = \tilde{w} \bar{x} + (1 - \tilde{w}) \mu_0,$$

where


$$\tilde{w} = \frac{\frac{n}{\tilde{\sigma}_n^2}}{\frac{n}{\tilde{\sigma}_n^2} + \frac{1}{\tau_0^2}}$$

# Normal example from Murphy ( $\lambda = 1/\sigma^2$ )


Initial values



After updating  $q_\mu$



After updating  $q_{\sigma^2}$



At convergence

